

Lecture #23: Code Generation, Part II

Parameter Passing Semantics: Value vs. Reference

- So far, our examples have dealt only with *value parameters*, which are the only kind found in C, Java, and Python

Ignorant comments from numerous textbook authors, bloggers, and slovenly hackers notwithstanding [End Rant].

- Pushing a parameter's value on the stack creates a copy that essentially acts as a local variable of the called function.
- C++ (and Pascal) have *reference parameters*, where assignments to the formal are assignments to the actual.

```
void incr(int& x) {                y = 4;
    x += 1;                        incr(y); // Now y == 5.
}
```

- The distinction is clear from this fact in Java:

```
var temp = x;    /* x a local variable. */
f(x);
/* At this point, x == temp, regardless of the body of f. */
```

Implementation of Reference Parameters

- Implementation of reference parameters is simple:
 - Push the address of the argument, not its value, and
 - To fetch from or store to the parameter, do an extra indirection.

```
void incr(int& x) {  
    x += 1;  
}
```

```
y = 4;  
incr(y);
```

```
incr:  
    # Prologue goes here  
    lw t0, 0(fp)  
    lw t1, 0(t0)  
    addi t1, t1, 1  
    sw t1, 0(t0)  
    # Epilogue goes here
```

```
# Assume y at -12(fp)  
li t0, 4  
sw t0, -12(fp)  
addi t0, fp, -12    # &y  
addi sp, sp, -4  
sw t0, 0(sp)  
jal incr  
addi sp, sp, 4
```

Copy-in, Copy-out Parameters

- Some languages, such as Fortran and Ada, have a variation on this: *copy-in, copy-out*. As for call by value, the value of the parameter is *copied into* the parameter, but also the final value of the parameter is *copied out* to the original location of the actual parameter after function returns.
 - “Original location” because of cases like $f(A[k])$, where k might change during execution of f . In that case, we want the final value of the parameter copied back to $A[k_0]$, where k_0 is the original value of k before the call.
 - Question: can you give an example where call by reference and copy-in, copy-out give different results?

Implementation of Copy-in/Copy-out Parameters

- We can implement copy-in/copy-out as a variation of the by-reference implementation.

```
void incr(int& x) {  
    x += 1; etc.  
}
```

```
y = 4;  
incr(y);
```

```
incr:  
    # Prologue goes here.  
    # Allocate local at -12(fp) for x  
    lw t0, 0(fp)  
    lw t0, 0(t0)  
    sw t0, -12(fp)    # Copy in  
    lw t0, -12(fp)  
    addi t0, t0, 1  
    sw t0, -12(fp)  
    # etc. (modify -12(fp) only)  
    lw t0, 0(fp)  
    lw t1, -12(fp)  
    sw t1, 0(t0)      # Copy out  
    # Epilogue goes here
```

```
# Assume y at -12(fp)  
li t0, 4  
sw t0, -12(fp)  
addi t0, fp, -12    # &y  
addi sp, sp, -4  
sw t0, 0(sp)  
jal incr  
addi sp, sp, 4
```

Parameter Passing Semantics: Call by Name

- Algol 60's definition says that the effect of a call $P(E)$ is as if the body of P were substituted for the call (dynamically, so that recursion works) and E were substituted for the corresponding formal parameter in the body (changing names to avoid clashes).
- It's a simple description that, for simple cases, is just like call by reference:

procedure F(x)	F(aVar);
integer x;	<i>becomes</i>
begin	aVar := 42;
x := 42;	
end F;	

- But the (unintended?) consequences were "interesting".

Call By Name: Jensen's Device

- Consider:

```
procedure DoIt (i, L, U, x, x0, E)
  integer i, L, U; real x, x0, E;
begin
  x := x0;
  for i := L step 1 until U do
    x := E;
  end DoIt;
```

- To set y to the sum of the values in array $A[1:N]$,

```
integer k;
DoIt(k, 1, N, y, 0.0, y+A[k]);
```

- To set z to the N th harmonic number:

```
DoIt(k, 1, N, z, 0.0, z+1.0/k);
```

- Now how are we going to make this work?

Call By Name: Implementation

- Basic idea: Convert call-by-name parameters into parameterless functions (traditionally called *thunks*.)
- To allow assignment, these functions can return the addresses of their results.
- So the call

```
DoIt(k, 1, N, y, 0.0, y+A[k]);
```

becomes something like (please pardon highly illegal notation):

```
integer t1;  real t2, t3, t4;  
t2 := 1.0; t3 := 0.0;  
DoIt(lambda: &k, lambda: &t2, lambda: &N, lambda: &y,  
      lambda: &t3, lambda: (t4 := y+A[k], &t4));
```

- Later languages have abandoned this particular parameter-passing mode.

One-dimensional Arrays

- How do we process retrieval from and assignment to $x[i]$, for an array x ?
- We assume that all items of the array have fixed size— S bytes—and are arranged sequentially in memory (the usual representation).
- Easy to see that the address of $x[i]$ must be

$$\&x + S \cdot i,$$

where $\&x$ is intended to denote the address of the beginning of x .

- Generically, we call such formulae for getting an element of a data structure *access algorithms*.
- The IL might look like this:

```
 $t_0 = \text{cgen}(\&A[E], t_0):$   
   $t_1 = \text{cgen}(\&A)$   
   $t_2 = \text{cgen}(E)$   
   $\Rightarrow t_3 := t_2 * S$   
   $\Rightarrow t_0 := t_1 + t_3$ 
```

Multi-dimensional Arrays

- A 2D array is a 1D array of 1D arrays.
- Java uses arrays of pointers to arrays for >1D arrays.
- But if row size constant, for faster access and compactness, may prefer to represent an $M \times N$ array as a 1D array of M 1D rows of length N (not pointers to rows): *row-major order*...
- Or, as in FORTRAN, a 1D array of N 1D columns of length M : *column-major order*.
- So apply the formula for 1D arrays repeatedly—first to compute the beginning of a row and then to compute the column within that row:

$$\&A[i][j] = \&A + i \cdot S \cdot N + j \cdot S$$

for an M -row by N -column array stored in row-major order.

- Where does this come from? Assuming S , again, is the size of an individual element, the size of a row of N elements will be $S \cdot N$.

IL for $M \times N$ 2D array

```
t = cgen(&e1[e2,e3]):  
  # Compute e1, e2, e3, and N:  
  t1 = cgen(e1);  
  t2 = cgen(e2);  
  t3 = cgen(e3)  
  t4 = cgen(N) # (N need not be constant)  
  ⇒ t5 := t4 * t2  
  ⇒ t6 := t5 + t3  
  ⇒ t7 := t6 * S  
  ⇒ t := t7 + t1  
return t
```

Array Descriptors

- Calculation of element address $\&e1[e2, e3]$ has the form

$$VO + S1 \times e2 + S2 \times e3$$

, where

- $VO(\&e1[0, 0])$ is the *virtual origin*.
 - $S1$ and $S2$ are *strides*.
 - All three of these are constant throughout the lifetime of the array (assuming arrays of constant size).
- Therefore, we can package these up into an *array descriptor*, which can be passed in lieu of a pointer to the array itself, as a kind of “*fat pointer*” to the array:

$\&e1[0][0]$	$S \times N$	S
--------------	--------------	-----

Array Descriptors (II)

- Assuming that `e1` now evaluates to the address of a 2D array descriptor, the IL code becomes:

```
t = cgen(&e1[e2,e3]):  
  t1 = cgen(e1);      # Yields a pointer to a descriptor.  
  t2 = cgen(e2;  
  t3 = cgen(e3)  
⇒ t4 := *t1;          # The V0  
⇒ t5 := *(t1+4)       # Stride #1  
⇒ t6 := *(t1+8)       # Stride #2  
⇒ t7 := t5 * t2  
⇒ t8 := t6 * t3  
⇒ t9 := t4 + t7  
⇒ t10:= t9 + t8
```

(Here, we assume 32-bit quantities. Adjust the constants appropriately for 64-bit pointers and/or integers.)

Array Descriptors (III)

- By judicious choice of descriptor values, can make the same formula work for different kinds of array.
- For example, if lower bounds of indices are 1 rather than 0, must compute address

$$\&e[1,1] + S1 \times (e2-1) + S2 \times (e3-1)$$

- But some algebra puts this into the form

$$V0' + S1 \times e2 + S2 \times e3$$

where

$$V0' = \&e[1,1] - S1 - S2 = \&e[0,0] \text{ (if it existed).}$$

- So with the descriptor

$V0'$	$S \times N$	S
-------	--------------	-----

we can use the same code as on the last slide.

- By passing descriptors as array parameters, we can have functions that adapt to many different array layouts automatically.

Other Uses for Descriptors

- No reason to stop with strides and virtual origins: can include other data.
- By adding upper and lower index bounds to a descriptor, can easily implement bounds checking.
- This also allows for runtime queries of array sizes and bounds.
- Descriptors also allow *views* of arrays: nothing prevents multiple descriptors from pointing to the same data.
- This allows effects such as slicing, array reversal, or array transposition without copying data.

Examples

- Consider a simple base array (in C):

```
int data[12] = { 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12 };
```

and descriptor types (including lengths):

```
struct Desc1 { int* V0, int S1, int len1 };
```

```
struct Desc2 { int* V0, int S1, int len1, int S2, int len2 };
```

- Here are some views:

```
Desc1 v0 = { data, 4, 12 }; /* All of data. */
```

```
Desc1 v1 = { &data[3], 4, 3 }; /* data[3:6]: [4, 5, 6]. */
```

```
/* Every other element of data: [1, 3, ...] */
```

```
Desc1 v2 = { data, 8, 6 };
```

```
Desc1 v3 = { &data[11], -4, 12 }; /* Reversed: [12, 11, ...] */
```

```
/* As a 2D 4x3 array: [ [ 1, 2, 3 ], [ 4, 5, 6 ], ... ] */
```

```
Desc2 v4 = { data, 12, 4, 4, 3 };
```

```
/* As row 2 of v4: [7, 8, 9] */
```

```
Desc1 v5 = { &data[6], 4, 3 }
```


Caveats

- Unfortunately, TANSTAAFL (There Ain't No Such Thing As A Free Lunch):
- Use of descriptors is nifty, but it costs:
 - For 1-D arrays, multiplication by a stride can be somewhat faster if the stride is known and is a power of 2 than when the stride is unknown due to difference in cost of multiplication vs. shift.
 - Fetching the VO from memory can also cost cycles relative to computing address of array on the stack or in static memory.
 - And fetching strides from memory is more expensive than using immediates.
 - Also, when stride is unknown can be hard to use vectorizing operations.