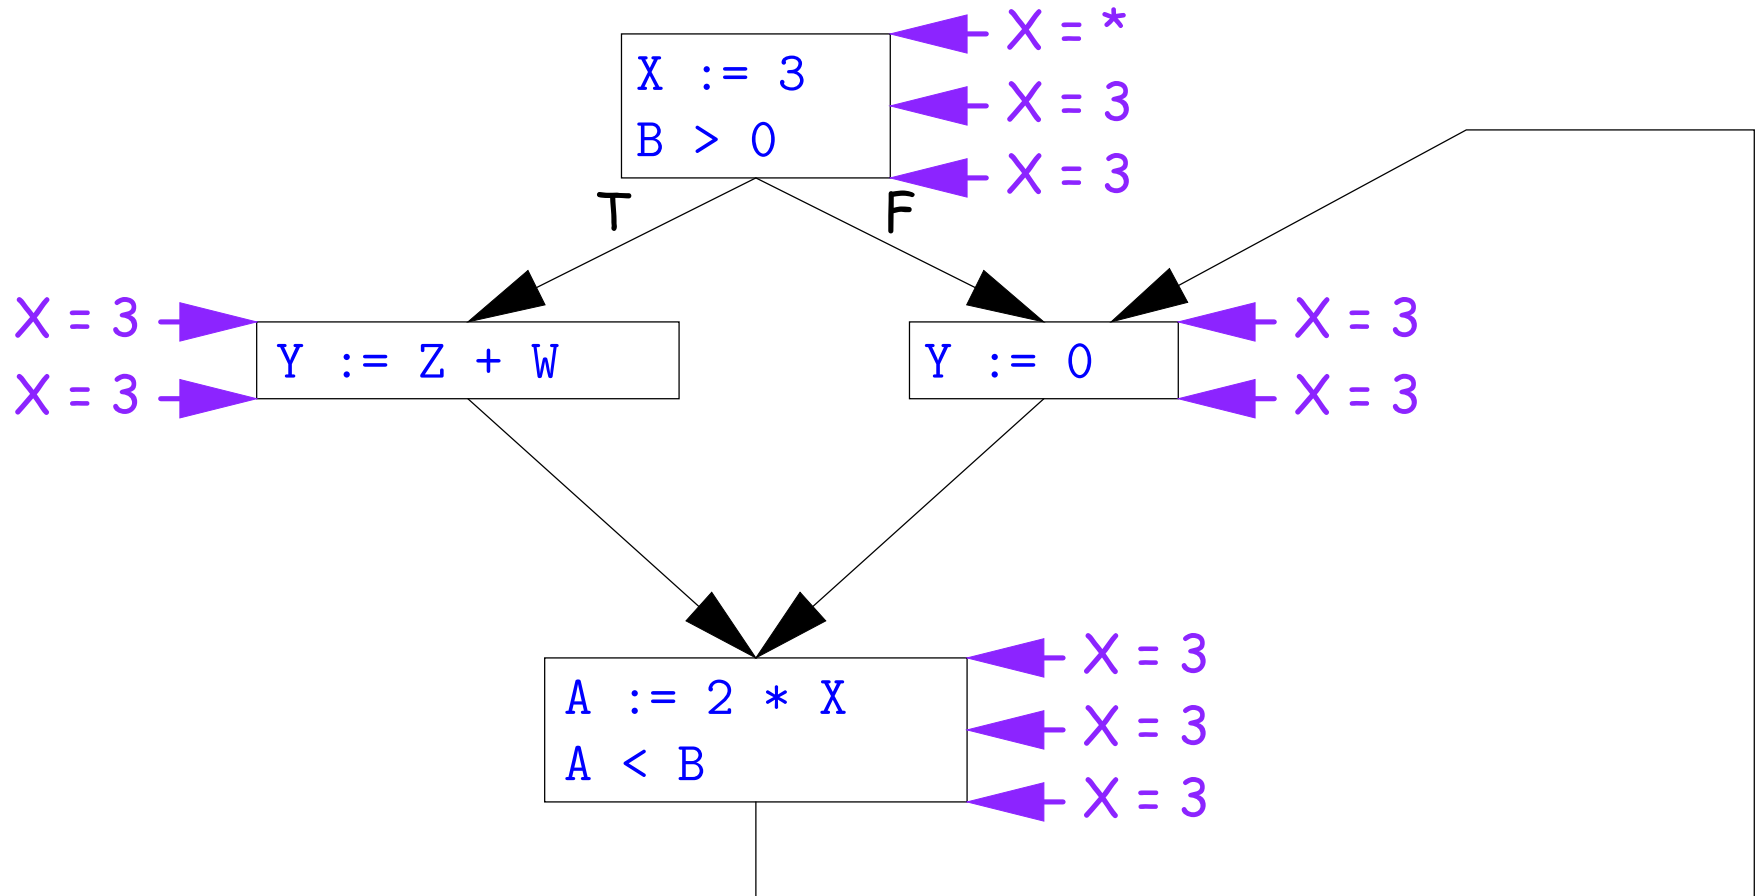


# Lecture 25: Code Optimization, Part 2

[Adapted in part from notes by R. Bodik and G. Necula]

# Liveness Analysis

Once constants have been globally propagated, we would like to eliminate dead code



After constant propagation, `X := 3` is dead code (assuming this is the entire CFG)

# Terminology: Live and Dead

- In the program


`X := 3; /*(1)*/ X = 4; /*(2)*/ Y := X /*(3)*/`

- the variable `X` is *dead* (never used) at point (1), *live* at point (2), and may or may not be live at point (3), depending on the rest of the program.
- More generally, a variable `x` is live at statement `s` if
  - There exists a statement `s'` that uses `x`;
  - There is a path from `s` to `s'`; and
  - That path has no intervening assignment to `x`
- A statement `x := ...` is dead code (and may be deleted) if `x` is dead after the assignment.

# Computing Liveness

- We can express liveness as a function of information transferred between adjacent statements, just as in copy propagation
- Liveness is simpler than constant propagation, since it is a boolean property (true or false).
- That is, the lattice has two values, with `false < true`.
- It also differs in that liveness depends on what comes *after* a statement, not before—we propagate information *backwards* through the flow graph, from `Lout` (liveness information at the end of a statement) to `Lin`.

# Liveness Rule 1




- So

$\text{Lout}(x, p) = \text{lub} \{ \text{Lin}(x, s) \text{ such that } p \text{ is a predecessor of } s \}.$

- Here, least upper bound (**lub**) is the same as "or".


## Liveness Rule 2



$\text{Lin}(X, s) = \text{true}$  if  $s$  uses the previous value of  $X$ .


- The same rule applies to any other statement that uses the value of  $X$ , such as tests (e.g.,  $X < 0$ ).

## Liveness Rule 3



$\text{Lin}(X, X := e) = \text{false}$  if  $e$  does not use the previous value of  $X$ .

## Liveness Rule 4



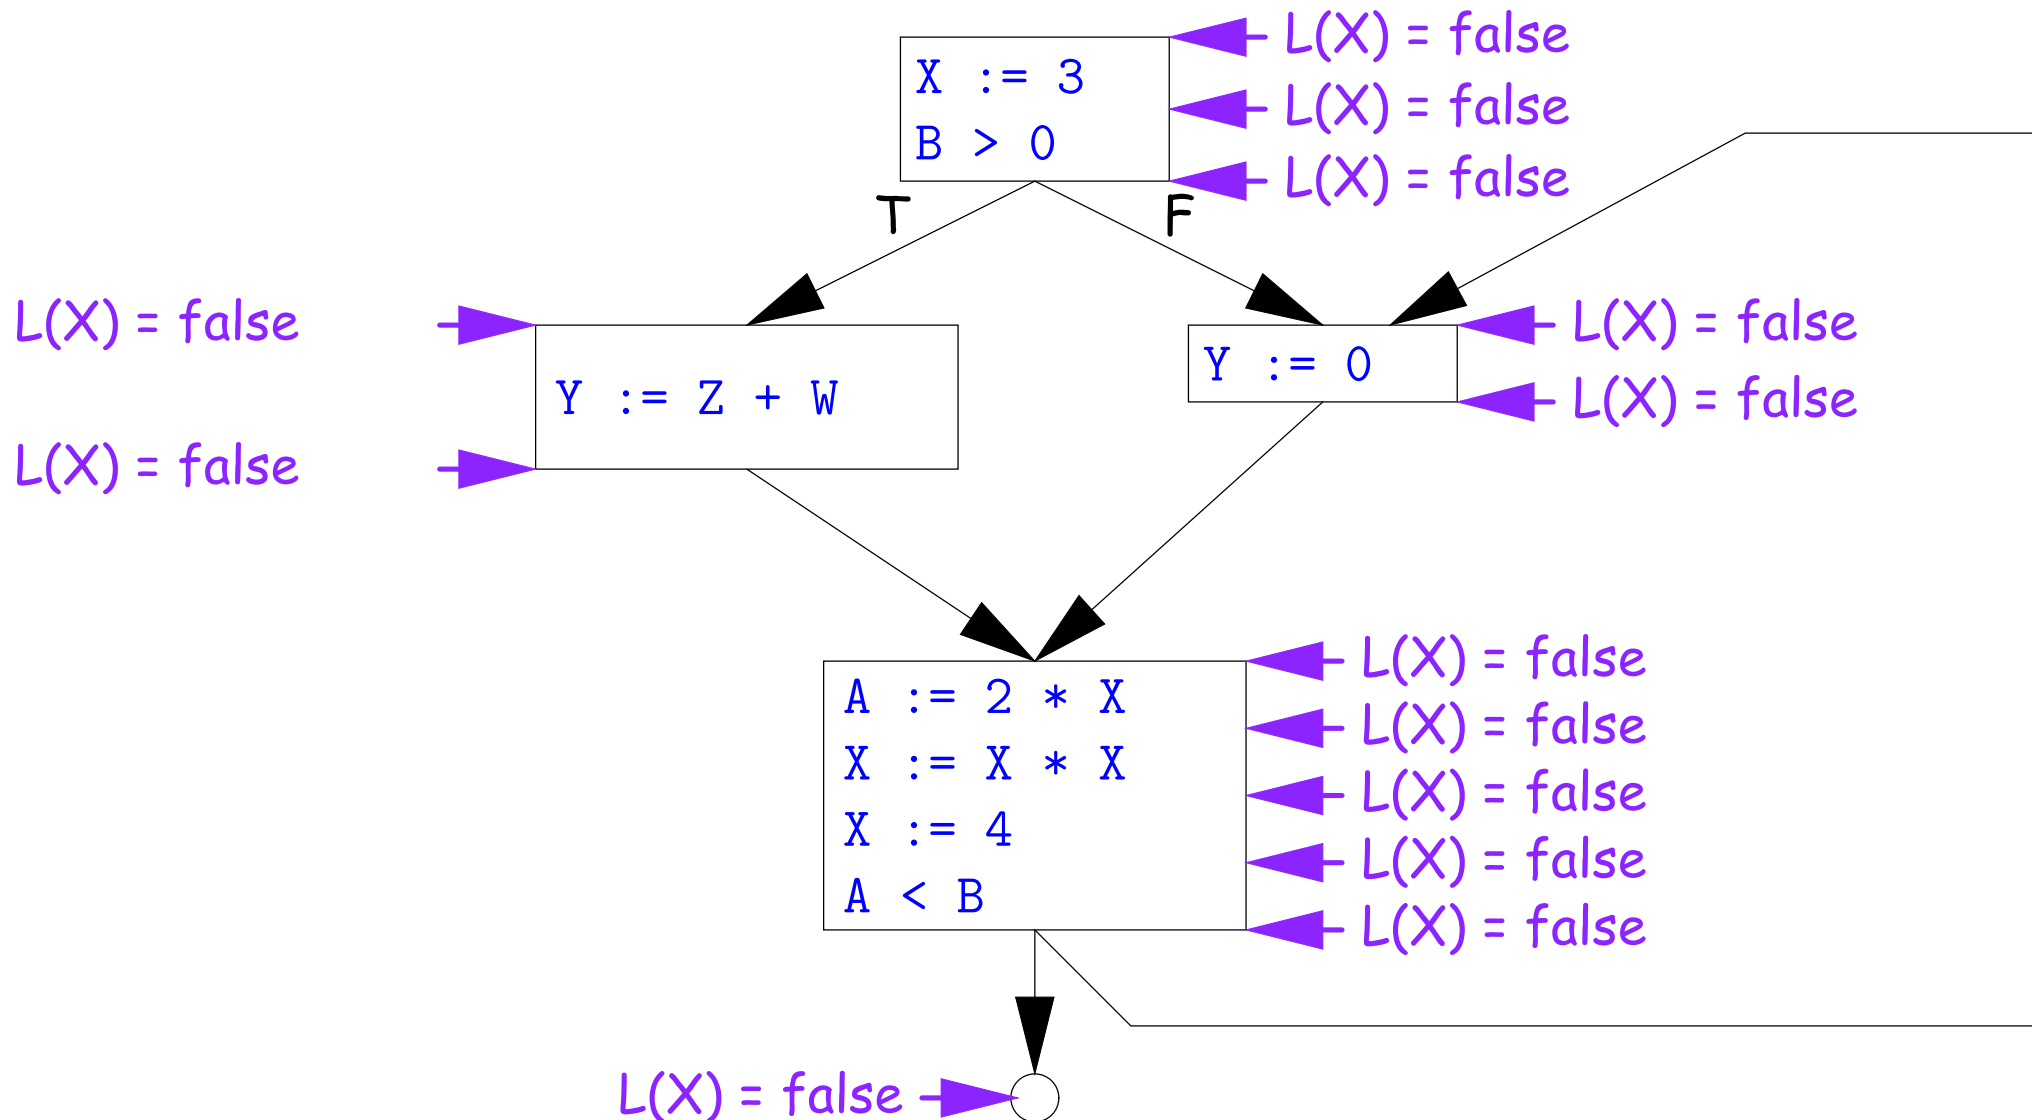
$L_{out}(X, s) = L_{in}(X, s)$  if  $s$  does not mention  $X$ .



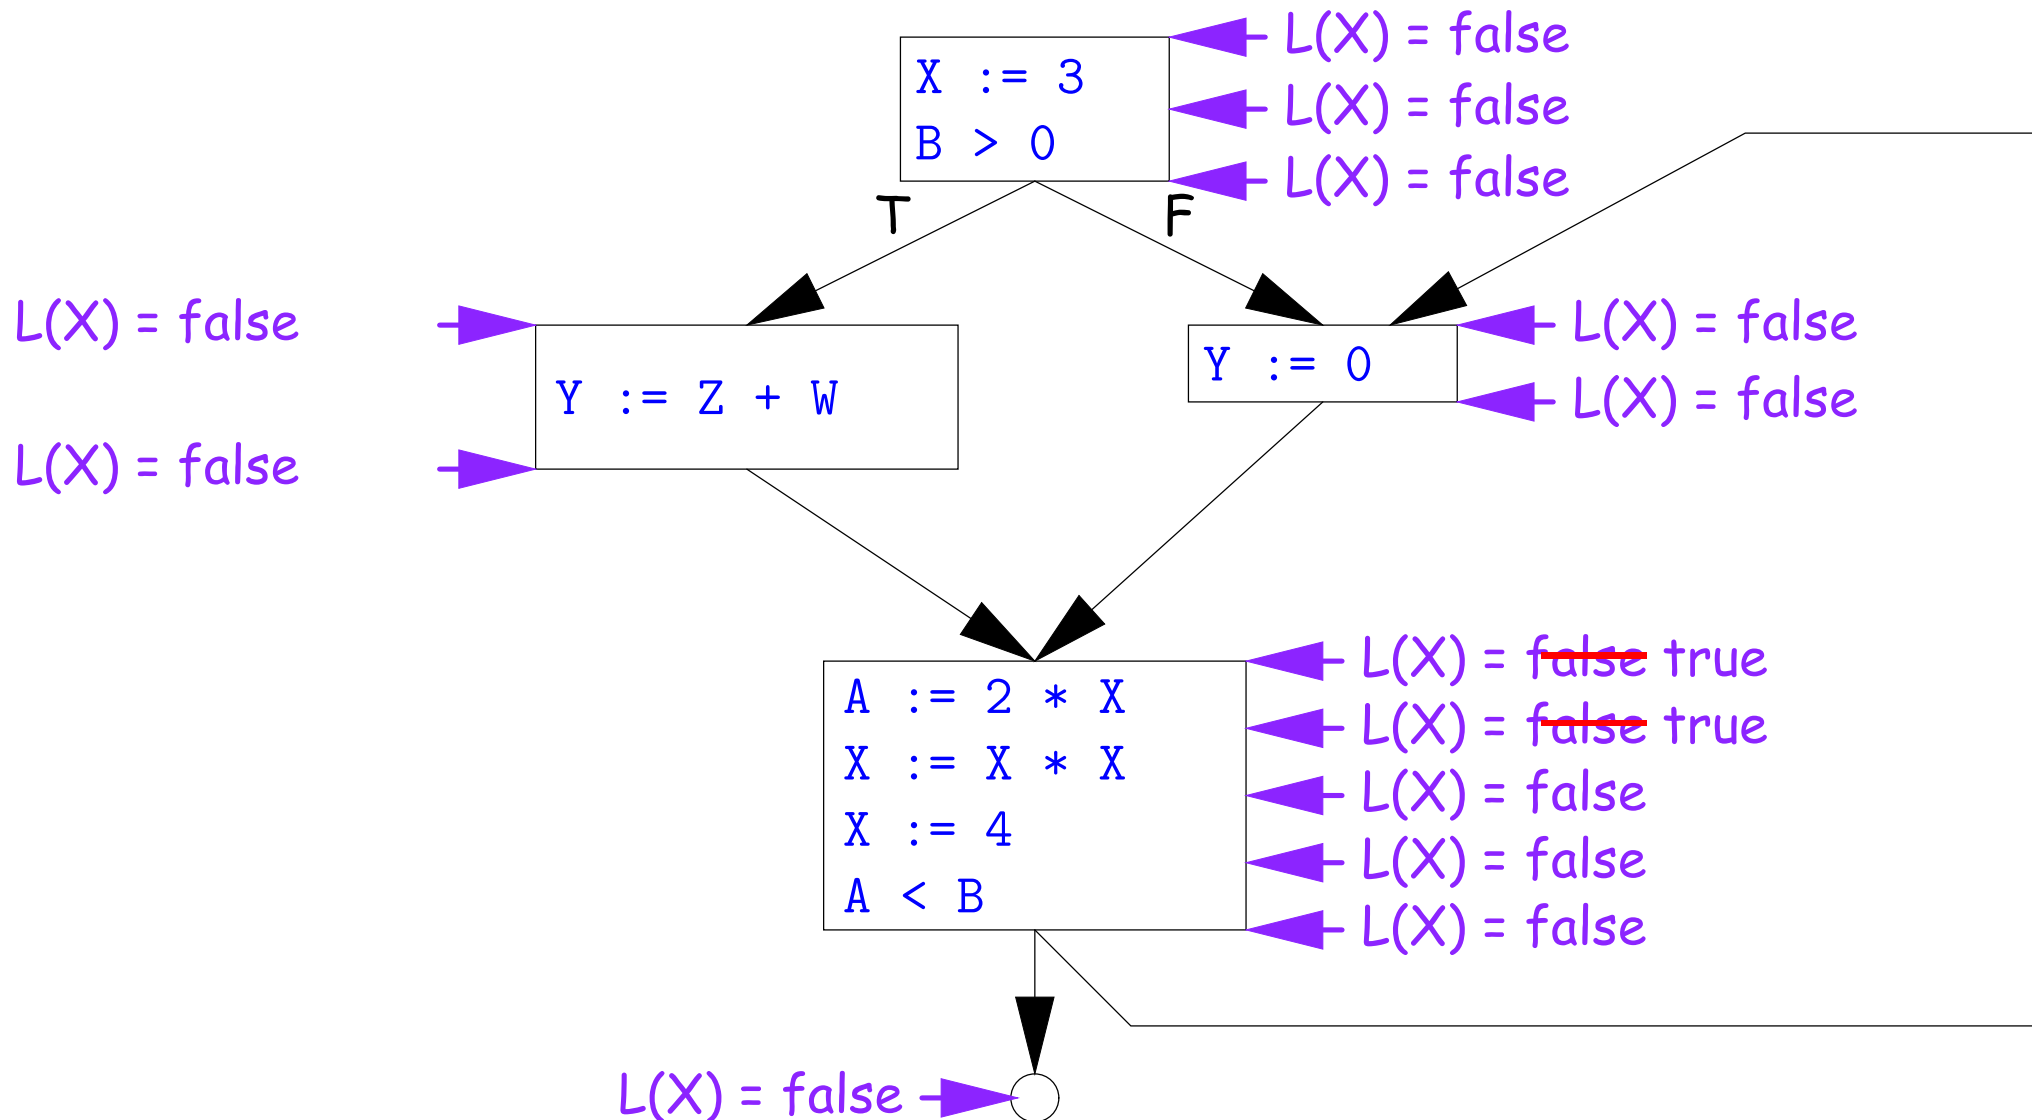
# Propagation Algorithm for Liveness

- Initially, let all **Lin** and **Lout** values be false.
- Set **Lout** value at the program exit to true iff **x** is going to be used elsewhere (e.g., if it is global and we are analyzing only one procedure).
- As before, repeatedly pick **s** where one of 1-4 does not hold and update using the appropriate rule, until there are no more violations.
- When we're done, we can eliminate assignments to **x** if **x** is dead at the point after the assignment.

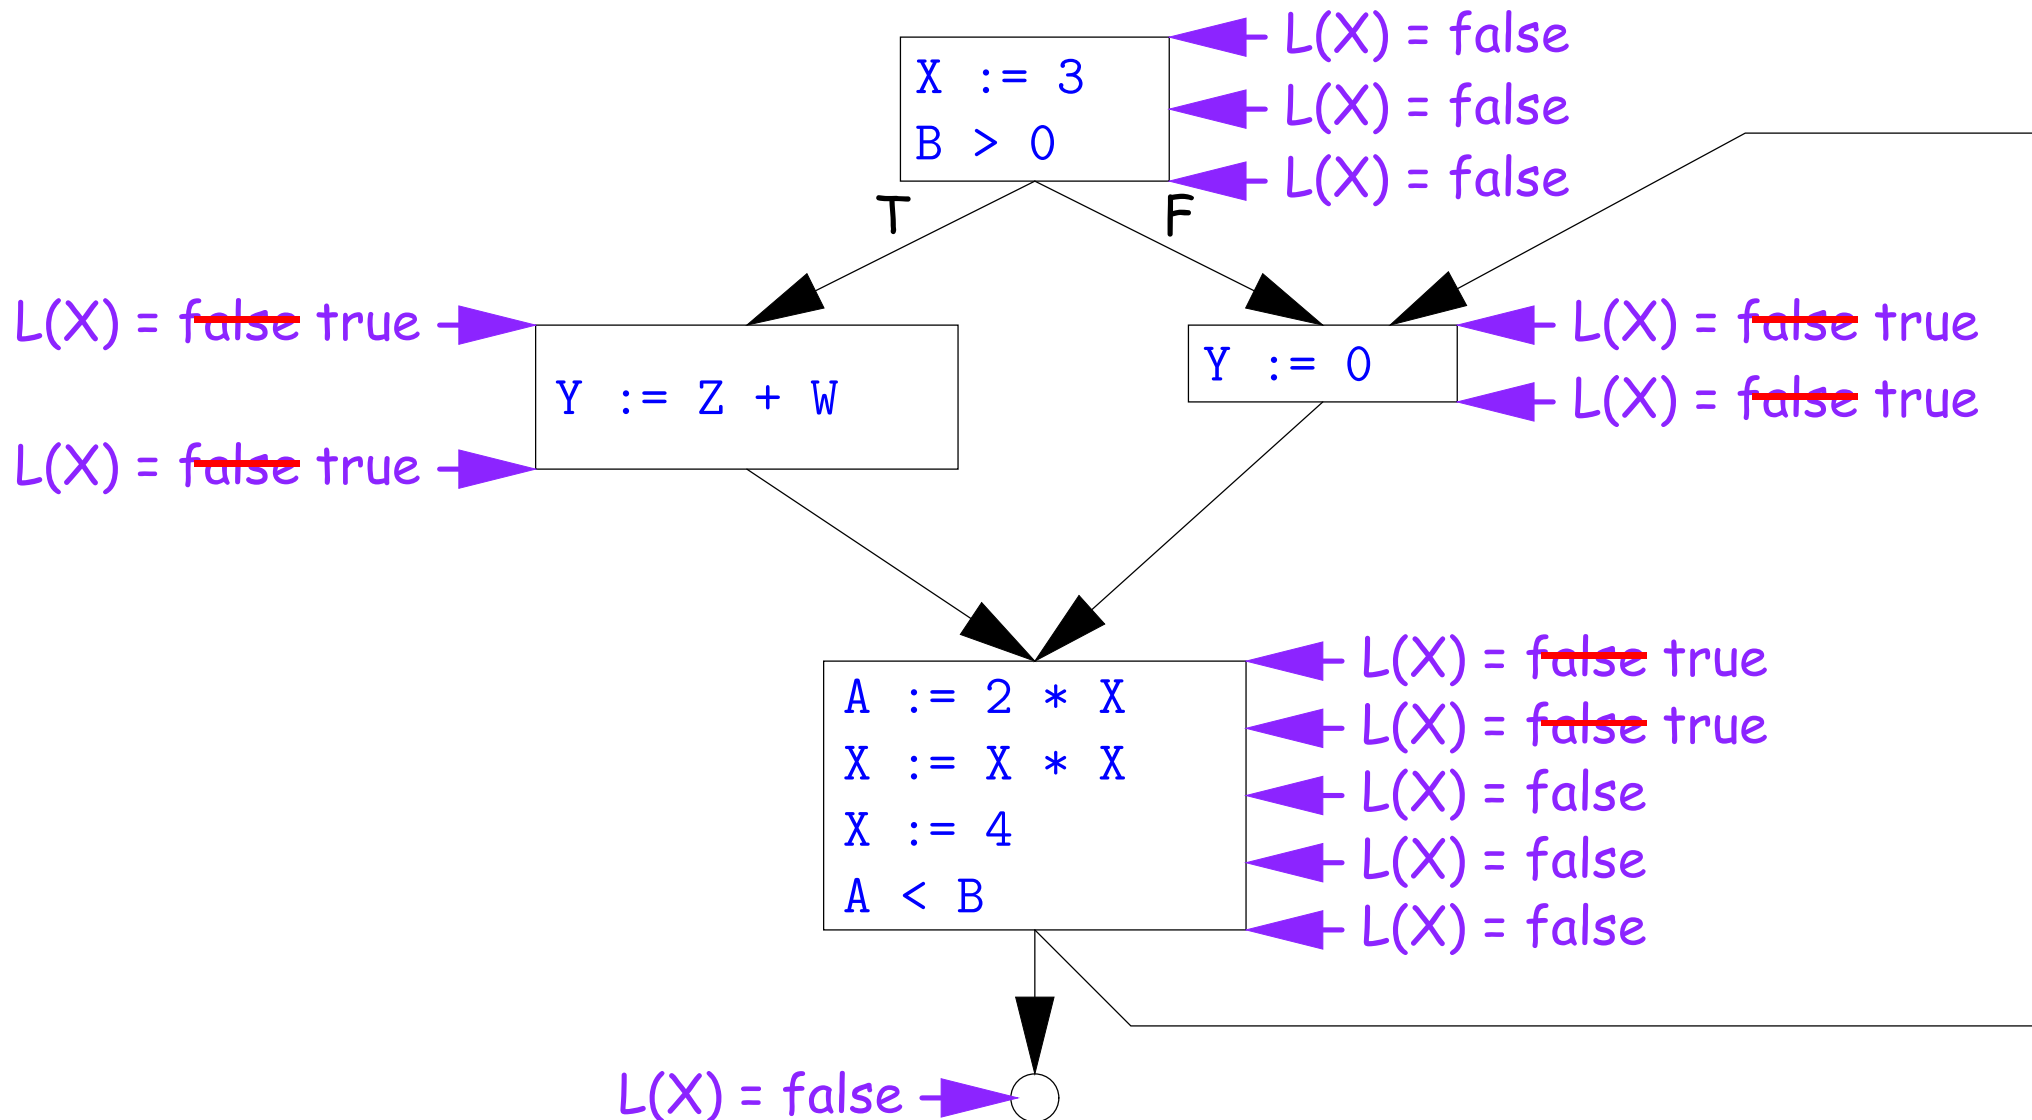
# Example of Liveness Computation



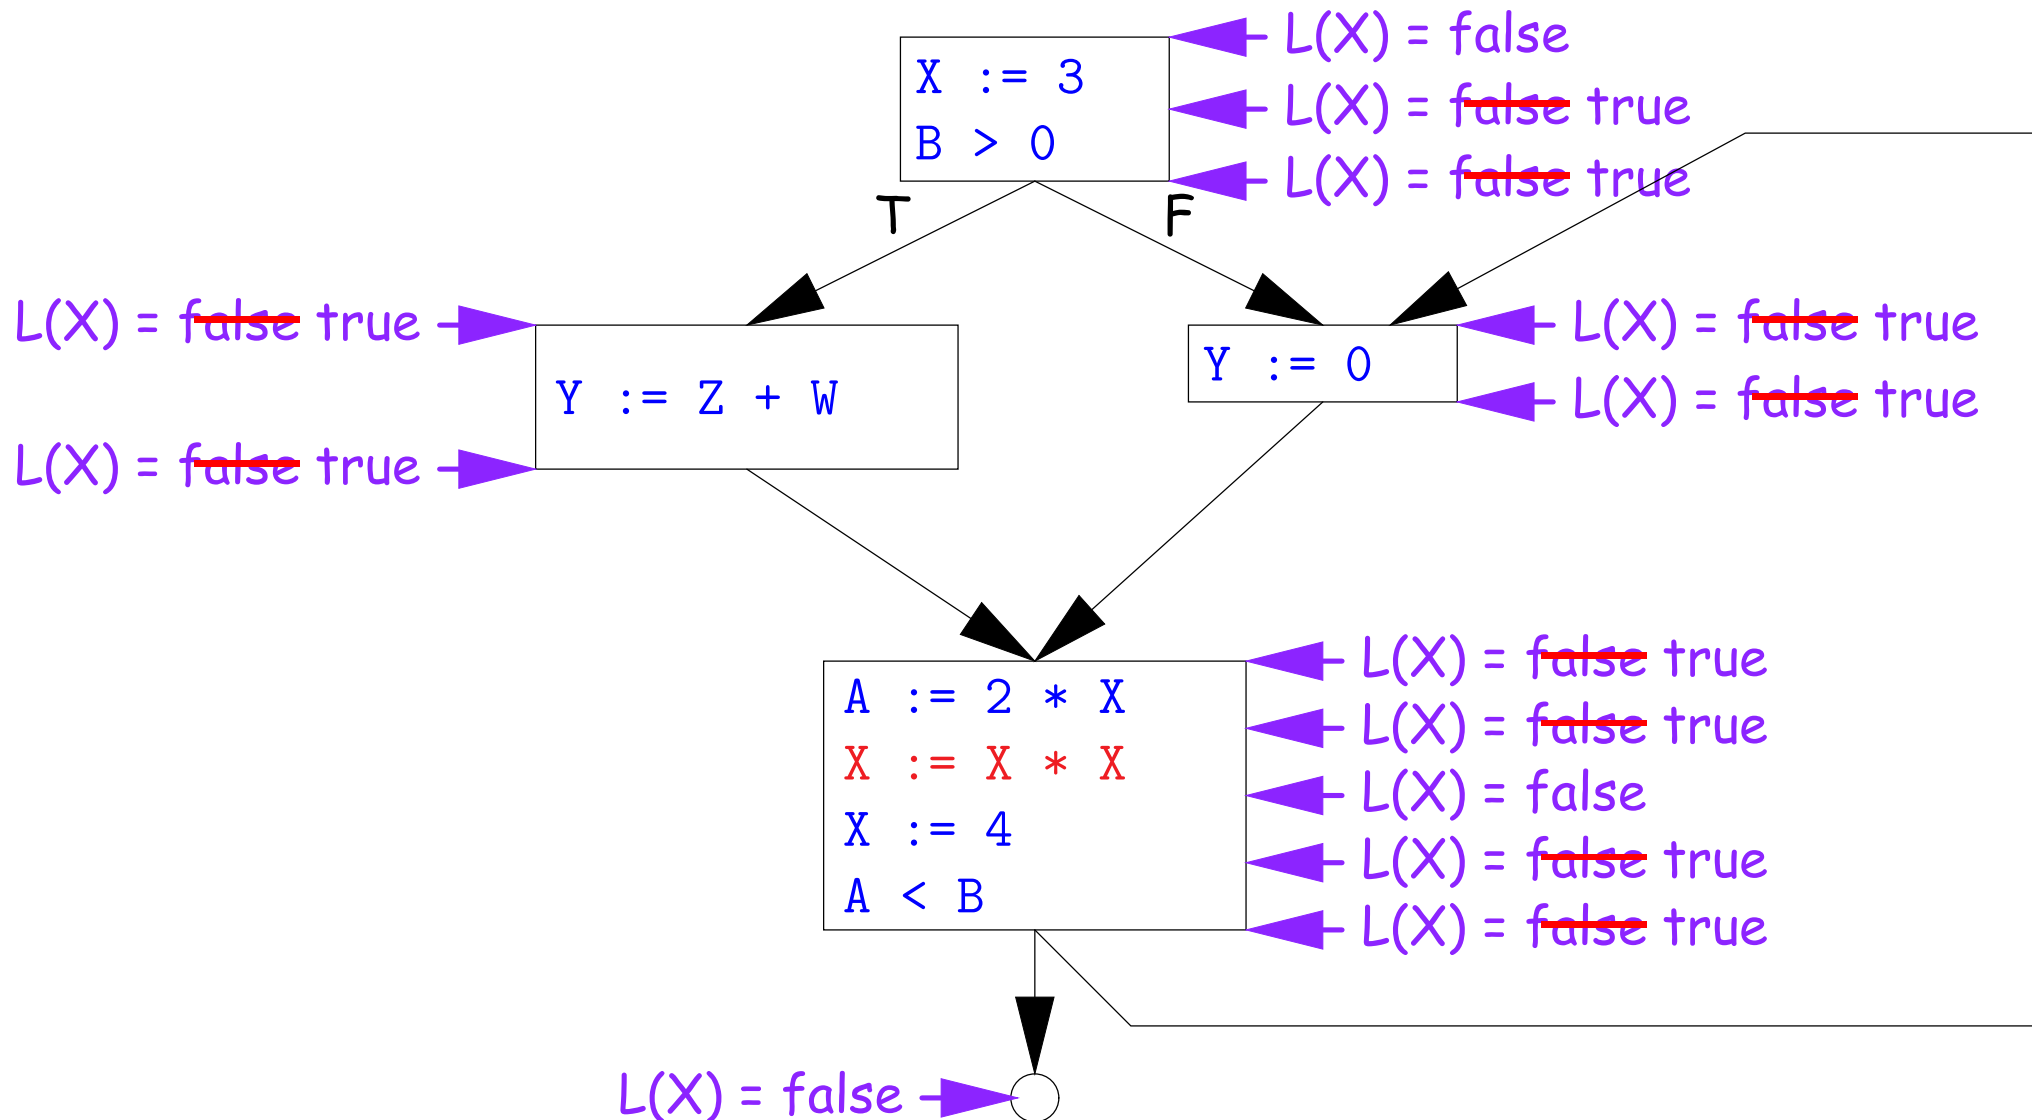
# Example of Liveness Computation



# Example of Liveness Computation



# Example of Liveness Computation



# Termination

- As before, a value can only change a bounded number of times: the bound being 1 in this case.
- Termination is guaranteed
- Once the analysis is computed, it is simple to eliminate dead code, but having done so, we must recompute the liveness information.

# SSA and Global Analysis

- For local optimizations, the single static assignment (SSA) form was useful.
- But applying it to a full CFG requires a trick.
- E.g., how do we avoid two assignments to the temporary holding  $x$  after this conditional?

```
if a > b:
    x = a
else:
    x = b
# where is x at this point?
```

- Answer: a small kludge known as  $\phi$  “functions”
- Turn the previous example into this:

```
if a > b:
    x1 = a
else:
    x2 = b
x3 =  $\phi$ (x1, x2)
```

## $\phi$ Functions

- An artificial device to allow SSA notation in CFGs.
- In a basic block, each variable is associated with one definition,
- $\phi$  functions in effect associate each variable with a *set of possible definitions*.
- In general, one tries to introduce them in strategic places so as to minimize the total number of  $\phi$ s.
- Although this device increases number of assignments in IL, register allocation can remove many by assigning related IL registers to the same real register.
- Their use enables us to extend such optimizations as CSE elimination in basic blocks to *Global CSE Elimination*.
- With SSA form, easy to tell (conservatively) if two IL assignments compute the same value: just see if they have the same right-hand side. The same variables indicate the same values.



# Loops

- In a CFG, a loop is simply a set of basic blocks,  $L$ , containing an entry block,  $e$ , such that
  - All paths from the entry node of the entire CFG to a block in  $L$  include  $e$ ;
  - All predecessors of a node in  $L$  are also in  $L$  (except for  $e$ , which must have a predecessor outside  $L$ ).
  - Every node in  $L$  has a path in  $L$  back to  $e$ .
- Here, for example,

```
j = i+1;  
while (j < N)  
    A[j] = A[j] / A[i]
```

The entry node contains the test  $j < n$  and the rest of the loop is the node containing the assignment to  $A[j]$ , which then loops back to the entry.

# Invariant Code Motion

- Consider the loop

```
while (i < N)
    A[i] = A[i] + j * x;
```

- Since  $j * x$  does not change in the loop, we can rewrite this as

```
tmp = j * x;
while (i < N)
    A[i] = A[i] + tmp;
```

- This is an example of *invariant code motion out of a loop*.
- What tells us that  $j*x$  does not change?
- We see that all assignments to  $j$  and  $x$  that apply at the point where the product is computed are outside the loop.
- And this we can get by observing where the assignments to the SSA-form for those variables are.

# Code Motion Caveat

- Code motion is not always appropriate.
- If the code to be moved, has side effects, or might cause an exception, could change the results.
- If the code is expensive, you will increase the time required for the program when the loop is not executed.
- Hence, you will see compilers rewrite loops like this:

```
if (i < N) {  
    /* Preheader */  
    while (i < N)  
        A[i] = A[i] + j * x;  
}
```

where `Preheader` marks a spot where the compiler can insert a new block to hold code moved out of the loop.

# Summary

- We've seen two kinds of analysis:
  - Constant propagation is a *forward analysis*: information is pushed from inputs to outputs.
  - Liveness is a *backwards analysis*: information is pushed from outputs back towards inputs.
- But both make use of essentially the same algorithm.
- Numerous other analyses fall into these categories, and allow us to use a similar formulation:
  - An abstract domain (abstract relative to actual values);
  - Local rules relating information between consecutive program points around a single statement; and
  - Lattice operations like least upper bound (or *join*) or greatest lower bound (or *meet*) to relate inputs and outputs of adjoining statements.

# Register Allocation

- Memory Hierarchy Management
- Register Allocation:
  - Register interference graph
  - Graph coloring heuristics
  - Spilling
- Cache Management

# The Memory Hierarchy

Computers employ a variety of memory devices, trading off capacity, persistence, and speed (some years ago):

Device	Access time (latency)	Capacity
Registers	1 cycle	256-2000 bytes
Cache	2-5 cycles	256KB-16MB
Main memory	100 cycles	32MB — >16GB
Disk	20K-10M cycles	10GB — > 1TB

# Managing the Memory Hierarchy

- Programs are written as if there are only two kinds of memory: main memory and disk (variables and files).
- Programmer is responsible for moving data from disk to memory.
- Hardware is responsible for moving data between memory and caches
- Compiler is responsible for moving data between memory and registers (which the programmer usually doesn't see).
- Cache and register sizes are growing slowly: important to manage them well.
- The cost of a cache miss is growing, and the widening gap is bridged with more caches.

# The Register Allocation Problem

- Our three-address code style uses temporaries profligately, simplifying code generation and optimization, but complicating final translation to assembly
- Hence, the register allocation problem:
  - Rewrite the intermediate code to use fewer temporaries than there are machine registers
- So we must assign more temporaries to a register, without changing the program behavior



# An Example

Consider the program

```
a := c + d
e := a + b
f := e - 1
```

assuming that assumption that *a* and *e* die after use. Then,

- Can reuse *a* after *a + b*
- Same with temporary *e* after *e - 1*
- Can allocate *a*, *e*, and *f* all to one register (*r1*):

```
r1 := c + d
r1 := r1 + b
r1 := r1 - 1
```

# Basic Register Allocation Idea

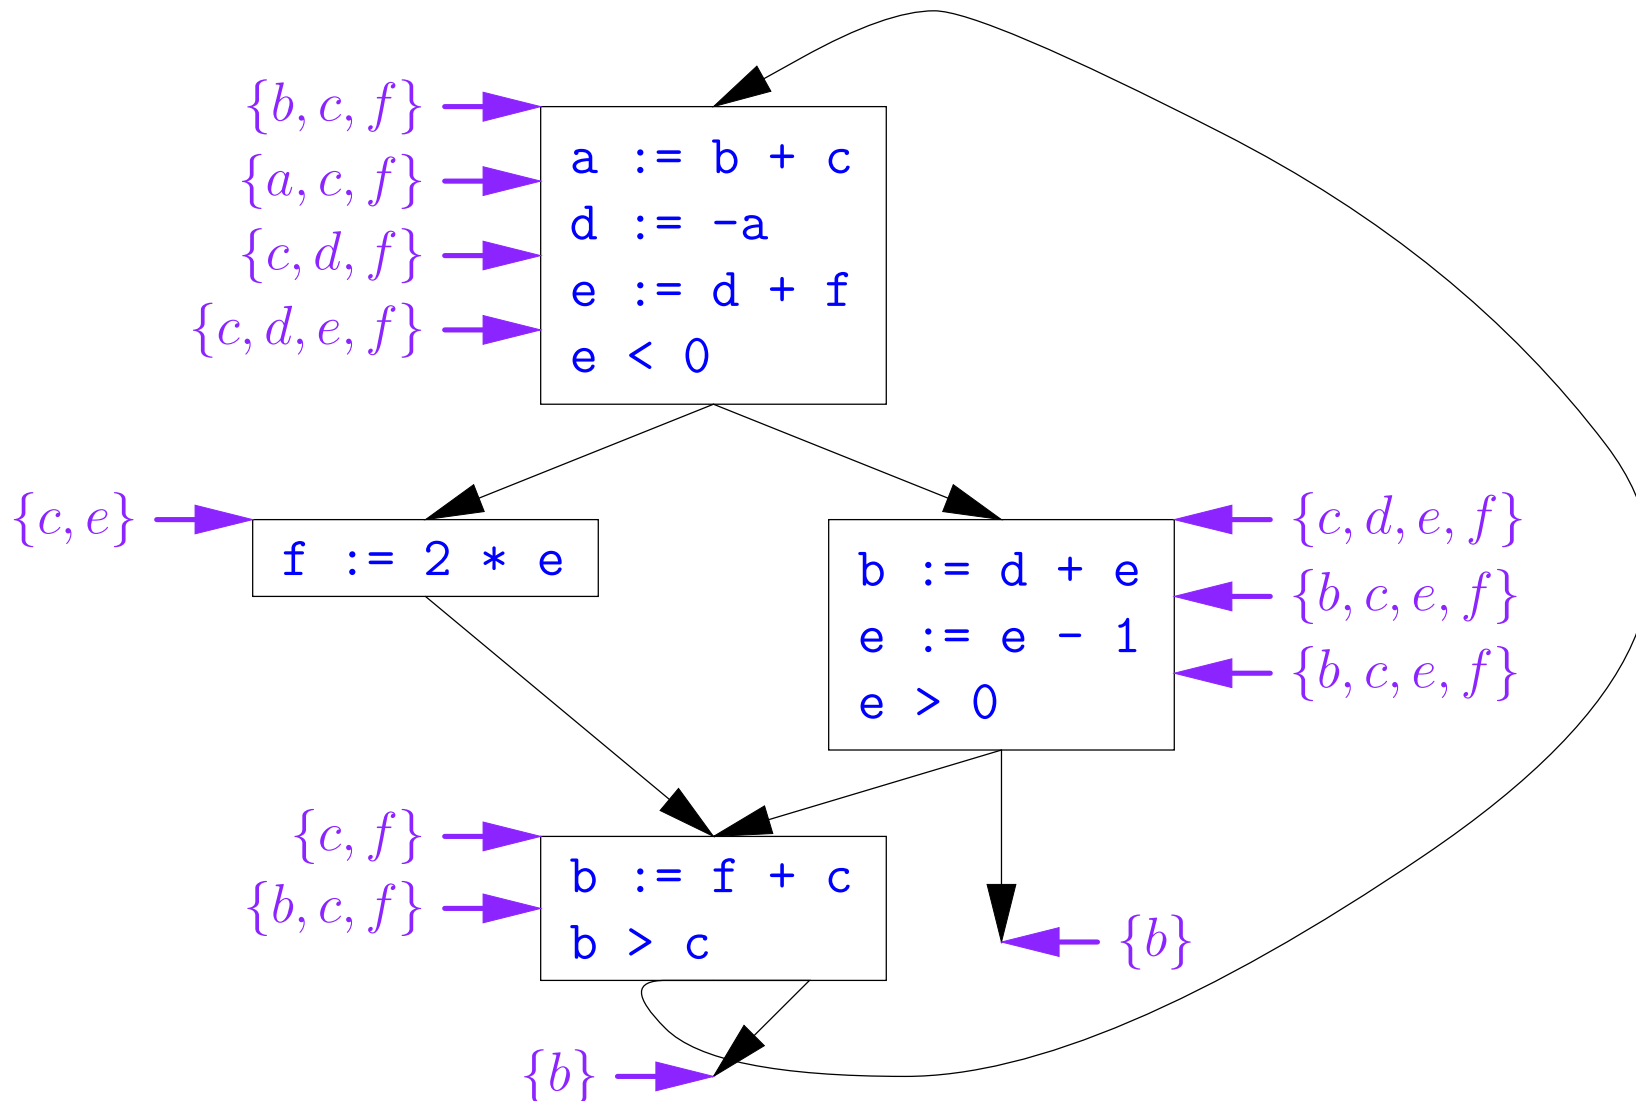
- So in general, since the value in a dead temporary is not needed for the rest of the computation,

Any set of temporaries can share a single physical register if at most one is alive at any program point.

- This rule is easy to apply to basic blocks. General CFGs are considerably trickier.


# Going Global: Allocation in CFGs (I)

First step is to compute live variables before each statement. In this example, assume that variable **b** is live at exit.



## Allocation in CFGs (II): Register Interference Graphs

- The sets in the previous slide indicate sets of virtual registers that are simultaneously alive at all points in the program, and therefore cannot share a physical register.
- Can summarize all these sets by constructing an undirected graph with a node for each virtual register, and an edge between any two virtual registers that appear together in the same set somewhere in the program.
- Call this the *register interference graph (RIG)*.



- The RIG extracts exactly the information needed to characterize legal register assignments
- Gives global (over the entire CFG) picture of the register requirements


## Allocation in CFGs (III): Graph Coloring

- A *coloring* of a graph is an assignment of colors to nodes, such that nodes connected by an edge have different colors.
- A graph is *k-colorable* if it has a coloring with *k* colors.
- In our problem, *colors = registers*. That is,

If we have *k* available machine registers and our register interference graph is *k*-colorable, then the coloring gives us a register assignment.

# Graph Coloring: Example

Consider the sample RIG:



- There is *no* coloring with fewer than 4 colors
- There *are* 4-colorings of this graph

Before

```
.
A: a := b + c
d := -a
e := d + f
if e >= 0 jump C
B: f := 2 * e
jump D
C: b := d + e
e := e - 1
if e <= 0 jump E
D: b := f + c
if b <= c jump A
E:
```

After

```
A: r2 := r3 + r4
r3 := -r2
r2 := r3 + r1
if r1 >= 0 jump C
B: r1 := 2 * r2
jump D
C: r3 := r3 + r2
r2 := r2 - 1
if r2 <= 0 jump E
D: r3 := r1 + r4
if r3 <= r4 jump A
E:
```

## Allocation in CFGs (III): Computing Graph Colorings

- The remaining problem is to compute a coloring for the interference graph.
- Unfortunately, this problem is hard (NP-hard). No guaranteed fast algorithms are known,
- And besides, a coloring might not exist for a given number of registers.
- For (1), we'll use heuristics.

# Graph Coloring Heuristic: Motivation

- Observation:
  - Pick a node  $t$  with  $< k$  neighbors in RIG.
  - Eliminate  $t$  and its edges from RIG.
  - If the resulting graph has a  $k$ -coloring then so does the original graph.
- Reason: whatever  $n \leq k - 1$  colors  $t$ 's neighbors have, we know we'll always be able to color  $t$  (since there are  $k$  colors). Therefore, eliminating  $t$  cannot affect the colorability of the other nodes.




# Graph Coloring Heuristic

The following works well in practice:

- Pick a node  $t$  with  $< k$  neighbors.
- Push  $t$  on a stack and remove it from the RIG.
- Repeat until the graph has no nodes.
- Then start popping nodes from the stack and adding them back to the graph, assigning colors to each as we go (starting with the last node added).
- At each step, we know we can pick a color different from those assigned to already colored neighbors, by the observation on the last slide.


# Example of Using the Heuristic (I)

Start with our sample RIG and with  $k = 4$ :



Stack: []

Now remove  $a$  and then  $d$ , giving



Stack: [ $d, a$ ] (top on left)

Now all nodes have  $< 4$  neighbors; remove. Stack is [ $f, e, b, c, d, a$ ].


## Graph Coloring Example (2)

- Now we assign colors ...er, ...registers to: *f, e, b, c, d, a* in that order.
- At each step, guaranteed there's a free register.

*r1 f* ●


## Graph Coloring Example (2)

- Now we assign colors ...er, ...registers to: *f, e, b, c, d, a* in that order.
- At each step, guaranteed there's a free register.




## Graph Coloring Example (2)

- Now we assign colors ...er, ...registers to:  $f, e, b, c, d, a$  in that order.
- At each step, guaranteed there's a free register.




## Graph Coloring Example (2)

- Now we assign colors ...er, ...registers to:  $f, e, b, c, d, a$  in that order.
- At each step, guaranteed there's a free register.




## Graph Coloring Example (2)

- Now we assign colors ...er, ...registers to:  $f, e, b, c, d, a$  in that order.
- At each step, guaranteed there's a free register.



## Graph Coloring Example (2)


- Now we assign colors ...er, ...registers to:  $f, e, b, c, d, a$  in that order.
- At each step, guaranteed there's a free register.





# Spilling


- What if during simplification we get to a state where all nodes have  $k$  or more neighbors?
- Example: try to find a 3-coloring of the RIG we've been using. After removing  $a$ , we get



- ...and now we are stuck, since all nodes have  $\geq 3$  neighbors.
- So, pick a node as a candidate for *spilling*, that is, to reside in memory.

## Example of Spilling


- Assume that  $f$  is picked as a candidate. When we remove it from the graph:



- Simplification now succeeds. We end up with the stack  $[e, c, b, d, f, a]$

## Example of Spilling (II)

- On the assignment phase we get to the point when we have to assign a color to  $f$
- Sometimes, it just happens that among the 4 neighbors of  $f$  we use  $< 3$  colors (*optimistic coloring*) ...



- ...but not this time.

## Example of Spilling (III)

- Since optimistic coloring failed we must spill register  $f$ : Allocate a memory location call it  $fa$  as the *home* of  $f$  (typically in the current stack frame).
- Before each operation that uses  $f$ , insert

$f := *fa$

- After each operation that defines (assigns to)  $f$ , insert

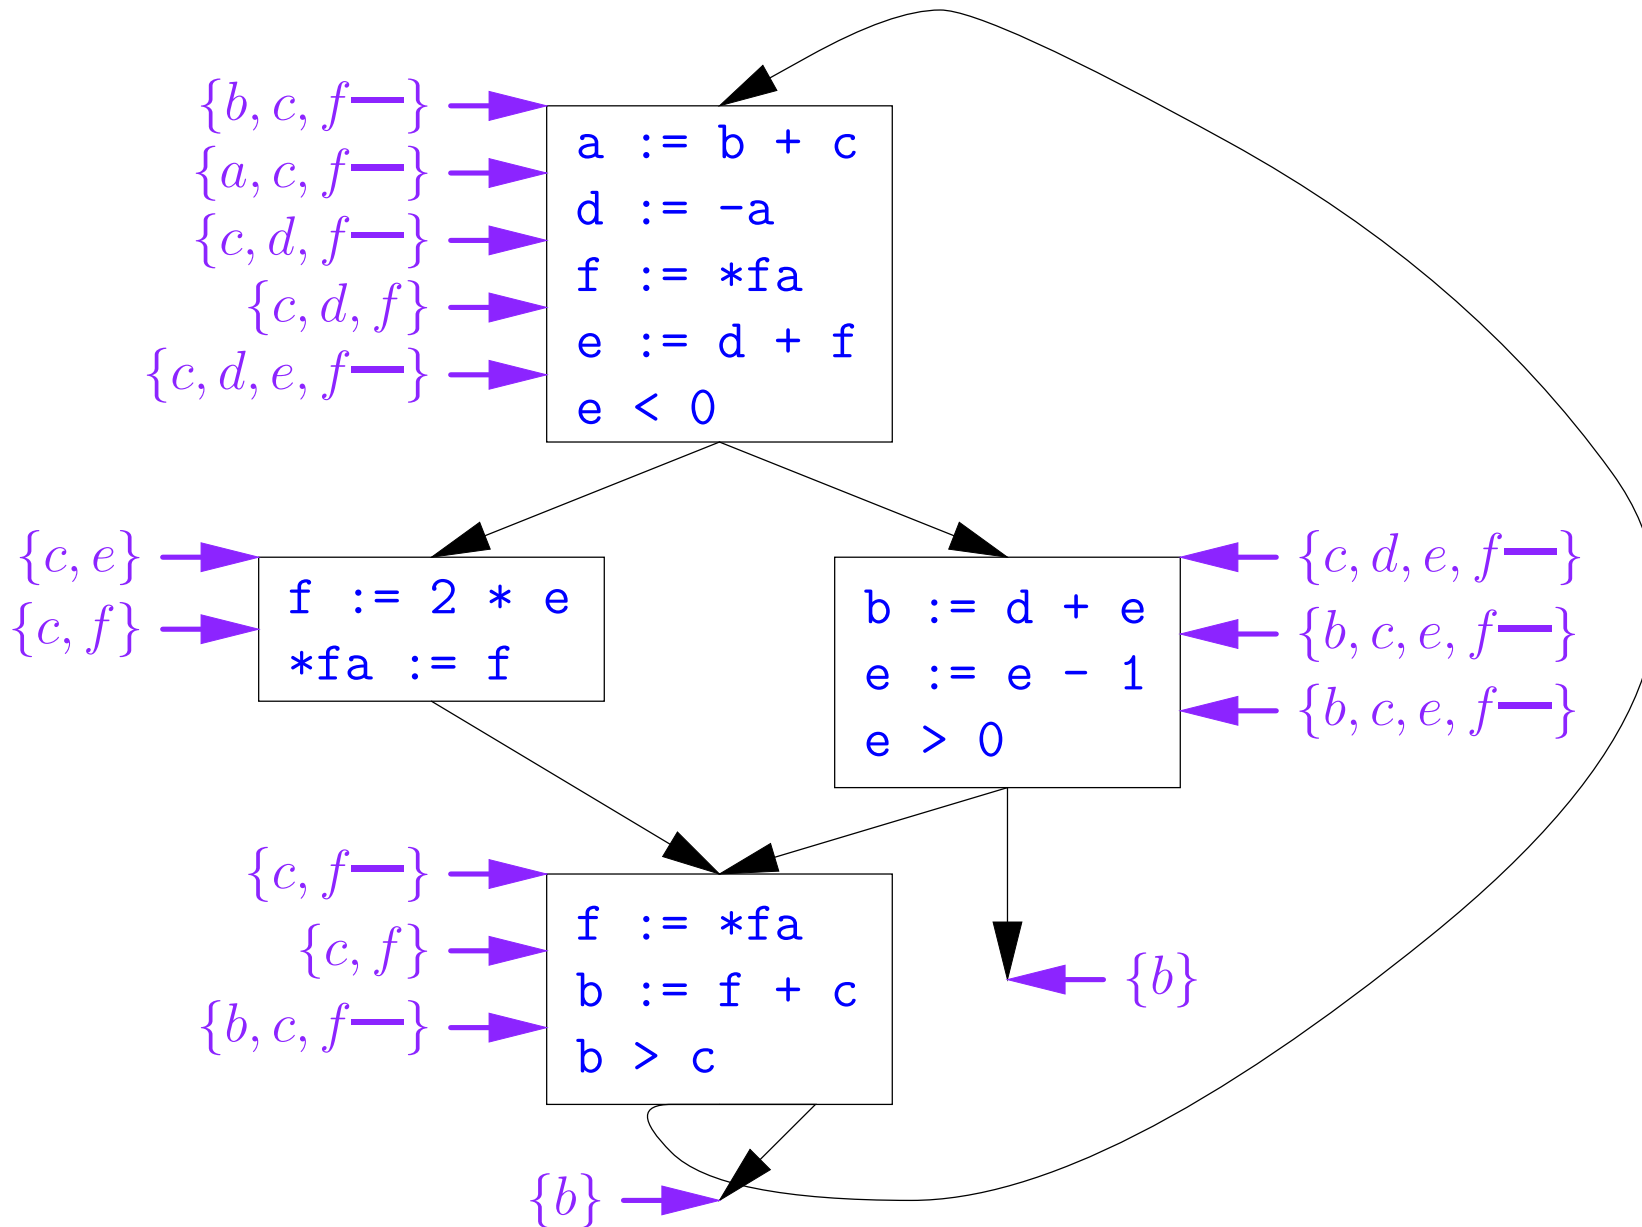
$*fa := f$

- This gives us:

```
A: a := b + c
d := -a
f := *fa
e := d + f
if e >= 0 jump C
B: f := 2 * e
*fa := f
jump D
```


```
C: b := d + e
e := e - 1
if e <= 0 jump E
f := *fa
D: b := f + c
if b <= c jump A
E:
```

# Recomputing Liveness Information



# A New RIG

- The new liveness information is almost as before, except that that  $f$  is live only
  - Between an  $f := *fa$  and the next instruction, and
  - Between a  $\text{store } f, fa$  and the preceding instruction.
- That is, spilling reduces the live range of  $f$ , and thus the registers it interferes with, giving us this RIG:



- And this graph is 3-colorable (left to the reader).

# What to Spill?

- In general, additional spills might be required to allow a coloring.
- The tricky part is deciding what to spill. Possible heuristics:
  - Spill temporaries with most conflicts
  - Spill temporaries with few definitions and uses
  - Avoid spilling in inner loops

# Caches

- Compilers are very good at managing registers (much better than programmers: the *C register* declaration is really obsolete).
- Caches are another matter. The problem is still left to programmers, and it is still an open question whether compilers can do much in general to improve performance
- But they can (and a few do) perform some simple cache optimization



# Cache Optimization

- Consider the loop

```
for(j = 1; j < 10; j += 1)
    for(i = 1; i < 1000000; i += 1)
        a[i] *= b[i]
```

- Why does this have terrible cache performance?
- On the other hand,

```
for(i = 1; i < 1000000; i += 1)
    for(j = 1; j < 10; j += 1)
        a[i] *= b[i]
```

computes the same thing, but with much better (possibly 10x) performance [again why?].

- Compilers can do this: *loop interchange*.

## Cache Optimization (II)

- Other kinds of memory layout decisions possible, such as *padding* rows of a matrix with extra bytes to avoid cache conflicts when traversing a column (or row in FORTRAN) of a matrix. [Why might that help?]
- *Prefetching* instructions on some hardware can inform cache of anticipated future memory fetches so that they can proceed in parallel. Again, it is possible for compilers to supply these to a limited extent.

# Summary

- Both because it eases code generation, greatly improves performance, and because it is difficult for programmers to do it for themselves, register allocation is a “must have” optimization in production compilers for standard procedural languages.
- Graph coloring is a powerful register allocation scheme that compilers can apply automatically
- Good cache management could give even larger payoffs, but so far is difficult.