

Lecture 27: Pointer Analysis, Storage Management

Today

- Points-to analysis: an instance of static analysis for understanding pointers [Based on slides from R. Bodik].
- Andersen's algorithm via deduction
- Storage Management and Garbage Collection

General Goals of Static Analysis

- Determine run-time properties statically at compilation.
- Sample property: "is variable x a constant?"
- Since we don't know the inputs, must consider all possible program executions.
- Conservative (err on the side of caution) for *soundness*:
 - allowed to say x is not a constant when it is,
 - but not that x is a constant when it is not .
- Many clients: optimization, verification, compilation.

Client 1: Optimizing virtual calls in Java

- Motivation: virtual calls are costly, due to method dispatch
- Idea:
 - determine the target of the call statically
 - if we can prove call has a single target method, call the target directly
- declared (static) types of pointer variables not precise enough for this, so, analyze their run-time (dynamic) types.

Client 1: Example

```
class A                { void foo() {...} }  
class B extends A     { void foo() {...} }  
void bar(A a)         { a.foo() } // OK to just call B.foo?  
B myB = new B();  
A myA = myB;  
bar(myA);
```

- Declared type of a permits a.foo() to target both A.foo and B.foo.
- Yet we know only B.foo is the target.
- What program property would reveal this fact?

Client 2: Verification of casts

- In Java, casts are checked at run time: `(Foo) e` translates to

```
if (! (e instanceof Foo))
    throw new ClassCastException()
```
- Java generics help readability, but still cast.
- The exception prevents any security holes, but is expensive.
- Static verification useful to catch bugs.
- Goal: prove that no exception will happen at runtime

Client 2: Example

```
class SimpleContainer { Object a;  
    void put (Object o) { a=o; }  
    Object get() { return a; } }  
SimpleContainer c1 = new SimpleContainer();  
SimpleContainer c2 = new SimpleContainer();  
c1.put(new Foo()); c2.put('Hello');  
Foo myFoo = (Foo) c1.get(); // Check not needed
```

What property will lead to desired verification?

Client 3: Non-overlapping fields in heap

```
E = new Thing (42);
for (j = 0; j < D.len; j += 1) {
    if (E.len >= E.max)) throw new OverflowException ();
    E.data[E.len] = D.data[i];  E.len += 1;
}
```

We assign to **E.len**, but we don't have to fetch from **D.len** every time; can save in register.

Pointer Analysis

- To serve these three clients, want to understand how pointers “flow,” that is, how they are copied from variable to variable.
- Interested in flow from *producers* of objects (*new Foo*) to *users* (*myFoo.f*).
- Complication: pointers may flow via the heap: a pointer may be stored in an object’s field and later be read from this field.
- For simplicity, assume we are analyzing Java without reflection, so that we know all fields of an object at compile time.

Analyses

- Client 1: virtual call optimization:
 - which producer expressions `new T()` produced the values that may flow to receiver `p` (a consumer) in a call?
 - Knowing producers tells us possible dynamic types of `p`, and thus also the set of target methods.
- Client 2: cast verification:
 - Same, but producers include expressions `(Type) p`.
- Client 3: non-overlapping fields: again, same question

Flow analysis as a constant propagation

- Initially, consider only new and assignments $p=r$:

```
if (...) p = new T1(); else p = new T2();  
r = p; r.f(); // what are possible dynamic types of r?
```

- We (conceptually) translate the program to

```
if (...) p =  $o_1$ ; else p =  $o_2$ ;  
r = p; r.f(); // what are possible symbolic constant values r?
```

Abstract objects

- The o_i constants are called abstract objects
- an abstract object o_i stands for any and all concrete objects allocated at the *allocation site* ('new' expression) with number i .
- When the analysis says a variable p may have value o_7 ,
- we know p may point to any object allocated at

`new7 Foo()`

Flow analysis: Add pointer dereferences

```
x = new Obj();      // o1
z = new Obj();      // o2
w = x;
y = x;
y.f = z;
v = w.f;
```

- To propagate the abstract objects through **p.f**, must keep track of the **heap state**—where the pointers point:
 - **y** and **w** point to same object
 - **z** and **y.f** point to same object, etc.

Flow-Insensitive Analysis


- The heap state may change at each statement, so ideally, track the heap state separately at each program point as in dataflow analysis.
- But to be scalable (i.e. practical), analyses typically don't do it.
- For example, to save space, can collapse all program points into one consequently, they keep a single heap state, and disregard the control flow of the program (*flow-insensitive* analysis):

assume that statements can execute in any order, and any number of times

- So, flow-insensitive analysis transforms this program

```
if (...) p = new T1(); else p = new T2();  
r = p; p = r.f;
```

into this CFG:



Flow-Insensitive Analysis, contd.

- Motivation: Just “version” of program state, hence less space
- Flow-insensitive analysis is *sound*, assuming we mean that *at least* all possible values of pointer from all possible executions found
- But it is generally *imprecise*:
 - In effect, adds many executions not present in the original program;
 - Does not distinguish value of p at various program points.

Canonical Statements

- Java pointers can be manipulated in complex statements, such as

`p.f().g.arr[i] = r.f.g(new Foo()).h`

- To keep complexity under control, prefer a small set of *canonical statements* that accounts for everything our analysis needs to serve as intermediate representation:

<code>p = new T()</code>	<code>new</code>
<code>p = r</code>	<code>assign</code>
<code>p = r.f</code>	<code>getfield</code>
<code>p.f = r</code>	<code>putfield</code>

- Complex statements can be canonicalized

`p.f.g = r.f` \implies `t1 = p.f; t2 = r.f; t1.g = t2`

- Can be done with a syntax-directed translation

Handling of method calls: Arguments and return values

- Translate calls into assignments. For example,

```
Object foo(T x) { return x.f }  
r = new T; s = foo(r.g)
```

could translate to

```
foo_retval = x.f;  
r = new T; x = r.g; s = foo_retval;
```

(have used flow-insensitivity: order irrelevant)

Handling of method calls: targets of virtual calls

- Call `p.f()` may call many possible methods
- To do the translation shown on previous slide, must determine what these targets are
- Suggest two simple methods:
 - Use declared type of `p`.
 - Check whole program to see which types are actually instantiated.

Handling of method calls: arrays

- We collapse all array elements into one.
- Represent this single element by a field named `arr`, so
`p.g[i] = r` becomes `p.g.arr = r`

Andersen's Algorithm for flow-insensitive points-to analysis

- Goal: computes a binary relation between variables and abstract objects:
 - o **flowsTo** x when abstract object o may be assigned to x .
- (Or, if you prefer, x **pointsTo** o .)
- Strategy: Deduce the flowsTo relation from program statements:
 - Statements are facts.
 - Analysis is a set of inference rules.
 - flowsTo relation is a set of facts inferred with analysis rules.

Statement facts

We'll write facts in the form x *predicate* y

$p = \text{new}_i T() \Rightarrow o_i \text{ new } p$

$p = r \Rightarrow r \text{ assign } p$

$p = r.f \Rightarrow r \text{ gf}(f) p$ (get field)

$p.f = r \Rightarrow r \text{ pf}(f) p$ (put field)

and apply these *inference rules*:

- Rule 1) $o_i \text{ new } p \Rightarrow o_i \text{ flowsTo } p$
- Rule 2) $o_i \text{ flowsTo } r \wedge r \text{ assign } p \Rightarrow o_i \text{ flowsTo } p$
- Rule 3) $o_i \text{ flowsTo } a \wedge a \text{ pf}(f) p \wedge p \text{ alias } r \wedge r \text{ gf}(f) b \Rightarrow o_i \text{ flowsTo } b$
- Rule 4) $o_i \text{ flowsTo } x \wedge o_i \text{ flowsTo } y \Rightarrow x \text{ alias } y$

Meaning of the results

- When the analysis infers *o flowsTo y*, what did we prove?
- Nothing useful, usually, since *o flowsTo y* does not imply that there is a program input for which *o* will *definitely* flow to *y*.
- BUT the useful results are places where analysis *does not* infer that *o flowsTo y*:
- In those cases—because the analysis assumes conservatively that *o* flows to *y* if there appears to be any possibility of that happening—we can infer that not *o flowsTo y* for all inputs.
- Same arguments apply to alias, pointsTo relations and many other static analyses in general.

Inference Example

The program:

```
x = new Foo(); // o1
z = new Bar(); // o2
w = x;
y = x;
y.f = z;
v = w.f;
```

The six facts:

```
o1 new x
o2 new z
x assign w
x assign y
z pf(f) y
w gf(f) v
```

Sample inferences:

```
o1 new x  $\Rightarrow$  o1 flowsTo x
o2 new z  $\Rightarrow$  o2 flowsTo z
o1 flowsTo x  $\wedge$  x assign w  $\Rightarrow$  o1 flowsTo w
o1 flowsTo x  $\wedge$  x assign y  $\Rightarrow$  o1 flowsTo y
o1 flowsTo y  $\wedge$  o1 flowsTo w  $\Rightarrow$  y alias w
o2 flowsTo z  $\wedge$  z pf(f) y  $\wedge$  y alias w  $\wedge$  w gf(f) v  $\Rightarrow$  o2 flowsTo v
etc.
```

Inference Example, contd.

- The inference must continue until no more facts can be derived; only then do we know we have performed sound analysis.
- In this example:
 - We have inferred o_2 *flowsTo* v
 - But we have *not* inferred o_1 *flowsTo* v .
 - Hence we know v will point only to instances of *Bar* (assuming the example contains the whole program)
 - Thus, casts *(Bar)* v will succeed
 - Similarly, calls $v.f()$ are optimizable.

New Topic: Storage Management and Garbage Collection

Storage Management





- Java has no means to free dynamic storage.
- However, when no expression in any thread can possibly be influenced by or change an object, it might as well not exist:

```
IntList wasteful ()
{
    IntList c = new IntList (3, new IntList (4, null));
    return c.tail;
    // variable c now deallocated, so no way
    // to get to first cell of list
}
```

- At this point, Java runtime, like Scheme's, recycles the object *c* pointed to: *garbage collection*.



Garbage Collection: Reference Counting

- Idea: Keep count of number of pointers to each object.



Garbage Collection: Mark and Sweep



Roots



- Start at roots (named variables, static and on stack)
- Perform graph traversal to find and **mark** all reachable storage.
- **Sweep** over memory, adding all unmarked storage to free list.

Garbage Collection: Mark and Sweep



Roots



- Start at roots (named variables, static and on stack)
- Perform graph traversal to find and **mark** all reachable storage.
- **Sweep** over memory, adding all unmarked storage to free list.

Garbage Collection: Mark and Sweep



Roots



- Start at roots (named variables, static and on stack)
- Perform graph traversal to find and **mark** all reachable storage.
- **Sweep** over memory, adding all unmarked storage to free list.

Garbage Collection: Mark and Sweep



Roots



- Start at roots (named variables, static and on stack)
- Perform graph traversal to find and **mark** all reachable storage.
- **Sweep** over memory, adding all unmarked storage to free list.

Garbage Collection: Mark and Sweep



Roots



- Start at roots (named variables, static and on stack)
- Perform graph traversal to find and **mark** all reachable storage.
- **Sweep** over memory, adding all unmarked storage to free list.

Garbage Collection: Mark and Sweep



Roots



- Start at roots (named variables, static and on stack)
- Perform graph traversal to find and **mark** all reachable storage.
- **Sweep** over memory, adding all unmarked storage to free list.

Garbage Collection: Mark and Sweep



Roots



- Start at roots (named variables, static and on stack)
- Perform graph traversal to find and **mark** all reachable storage.
- **Sweep** over memory, adding all unmarked storage to free list.

Garbage Collection: Mark and Sweep

Roots




- Start at roots (named variables, static and on stack)
- Perform graph traversal to find and **mark** all reachable storage.
- **Sweep** over memory, adding all unmarked storage to free list.


Copying Garbage Collection

- Copy (and move) only reachable (useful) storage from 'from' space to 'to' space.
- The 'from' and 'to' areas are called *semispaces*. Need twice the virtual memory you actually use.
- As you copy, mark 'from' storage as moved, and leave behind a *forwarding pointer* that tells how to translate other references to the old storage.
- At end of algorithm, 'from' and 'to' swap roles, and the old 'from' area is freed *en masse*.
- Copied storage is *compacted* (gaps squeezed out) with possible advantages for memory access.

Copying Garbage Collection, Illustrated




from:




to:

Roots





Copying Garbage Collection, Illustrated

Roots





----->
forwarding pointer

Roots



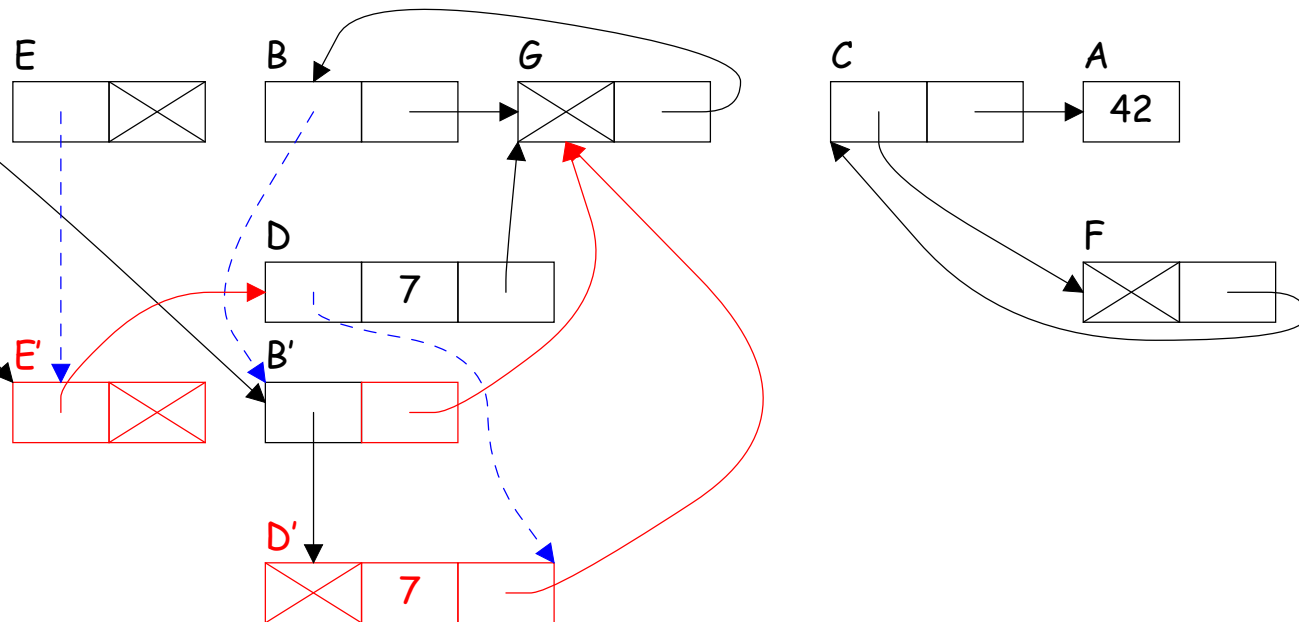

Copying Garbage Collection, Illustrated

Roots





----->
forwarding pointer

Roots





Copying Garbage Collection, Illustrated





Copying Garbage Collection, Illustrated

Roots





----->
forwarding pointer

Roots





Copying Garbage Collection, Illustrated

Roots





----->
forwarding pointer

Roots





Copying Garbage Collection, Illustrated

Roots




----->
forwarding pointer

Roots

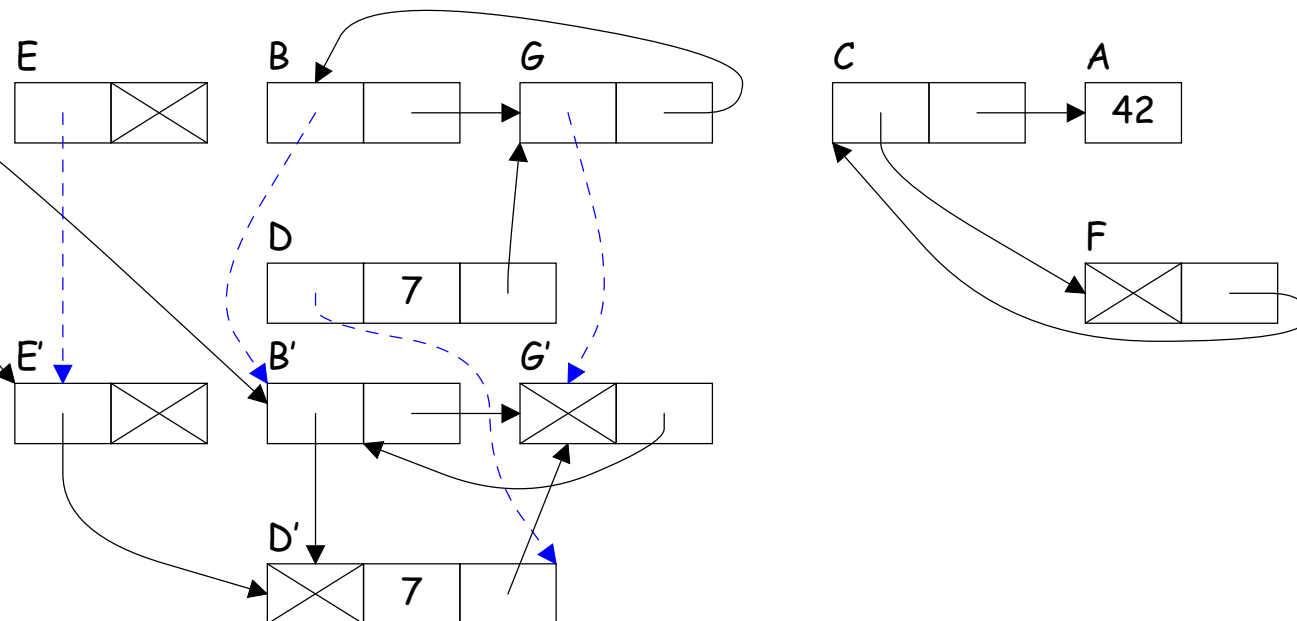


Copying Garbage Collection, Illustrated

Roots



Roots




Copying Garbage Collection, Illustrated


Roots



to:



Roots



Roots and Other Pointers

- Above methods require that we know locations of roots and of pointer fields in objects.
- Positions of some roots change during execution.
- Compiler keeps tables mapping PC to where roots are.
- Runtime type information (virtual tables) keep information of where pointer fields are.
- Implementation must guarantee that fields are initialized.

Conservative Garbage Collection

- With C, you have none of the needed information.
- But easy to know the addresses of allocated storage, and sizes of allocated objects (allocator keeps them around).
- So, **guess** that any word that looks like an address of allocated storage is a valid address.
- Do mark-and-sweep on this assumption (look at whole stack and static storage for roots).
- Marks some garbage, but can be surprisingly effective.

Generational Garbage Collection


- Heap storage tends to “die young.”
- So divide memory into young and old storage, and do copying only on young storage.
- Must add old storage that points to young storage to roots.
- When young storage survives a GC (or two), move it to old storage.
- Every now and then, stop the world and do a full garbage collection.
- This technique significantly speeds up GC.

Region-Based Allocation

- Garbage collection (all forms) does incur overheads, which can be unpredictable,
- While manual freeing is prone to error and inconvenient.
- One compromise is *region-based allocation*.
- Idea:
 - Create a data structure known as a region (or zone, or arena, or various other names).
 - Provides two operations: allocate object, and free *all* objects.
- Thus, to perform calculation that creates lots of temporary heap objects,
 - Create region (a local variable).
 - Allocate all the temporary storage in this region.
 - Delete whole region at end.


Region Implementation

- Simple implementation: allocate storage in big blocks, and allocate objects sequentially in the blocks.
- Freeing all blocks frees all the objects quickly.



Region Implementation


- Simple implementation: allocate storage in big blocks, and allocate objects sequentially in the blocks.
- Freeing all blocks frees all the objects quickly.



```
x = aRegion.alloc (40);
```

Region Implementation


- Simple implementation: allocate storage in big blocks, and allocate objects sequentially in the blocks.
- Freeing all blocks frees all the objects quickly.



```
x = aRegion.alloc (40);  
y = aRegion.alloc (100);
```

Region Implementation


- Simple implementation: allocate storage in big blocks, and allocate objects sequentially in the blocks.
- Freeing all blocks frees all the objects quickly.



```
x = aRegion.alloc (40);  
y = aRegion.alloc (100);  
z = aRegion.alloc (120);
```

Region Implementation


- Simple implementation: allocate storage in big blocks, and allocate objects sequentially in the blocks.
- Freeing all blocks frees all the objects quickly.



```
x = aRegion.alloc (40);  
y = aRegion.alloc (100);  
z = aRegion.alloc (120);  
v = aRegion.alloc (100);
```

Region Implementation

- Simple implementation: allocate storage in big blocks, and allocate objects sequentially in the blocks.
- Freeing all blocks frees all the objects quickly.



```
x = aRegion.alloc (40);  
y = aRegion.alloc (100);  
z = aRegion.alloc (120);  
v = aRegion.alloc (100);  
w = aRegion.alloc (50);
```

Region Implementation

- Simple implementation: allocate storage in big blocks, and allocate objects sequentially in the blocks.
- Freeing all blocks frees all the objects quickly.

aRegion: ☐

```
x = aRegion.alloc (40);  
y = aRegion.alloc (100);  
z = aRegion.alloc (120);  
v = aRegion.alloc (100);  
w = aRegion.alloc (50);  
aRegion.freeAll ();
```

Region Implementation

- Simple implementation: allocate storage in big blocks, and allocate objects sequentially in the blocks.
- Freeing all blocks frees all the objects quickly.

```
x = aRegion.alloc (40);  
y = aRegion.alloc (100);  
z = aRegion.alloc (120);  
v = aRegion.alloc (100);  
w = aRegion.alloc (50);  
aRegion.freeAll ();
```

- Potential problem: using x, y, z, ... after freeAll.